

CSCI4390/6390 – Data Mining
Fall 2009, Exam I
Total Points: 100 + (10 bonus)

1. (30 points) Let $\Sigma = \begin{pmatrix} 101/2 & 99/2 \\ 99/2 & 101/2 \end{pmatrix}$ be the covariance matrix for some dataset, with mean $\mu = (2, 5)$. Answer the following questions.

- (a) (10 points) Compute the dominant eigenvector and eigenvalue of Σ by the power method. Carry out at least 3 iterations, i.e., starting with an initial vector x_0 , iterate until you get x_3 . Approximate up to 2 decimal places, rounding up when necessary. Don't forget to normalize the eigenvector.

Answer: let $x_0 = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$.

We get $\Sigma \cdot x_0 = \begin{pmatrix} 50.5 \\ 49.5 \end{pmatrix} = \begin{pmatrix} 1 \\ 0.98 \end{pmatrix} = x_1$

Next $\Sigma \cdot x_1 = \begin{pmatrix} 99.01 \\ 98.99 \end{pmatrix} = \begin{pmatrix} 1 \\ 1 \end{pmatrix} = x_2$

Finally $\Sigma \cdot x_2 = \begin{pmatrix} 100 \\ 100 \end{pmatrix} = \begin{pmatrix} 1 \\ 1 \end{pmatrix} = x_3$

This implies $u_1 = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 \\ 1 \end{pmatrix}$ and $\lambda_1 = 100$.

- (b) (5 points) Using the spectral decomposition and/or using the fact that eigenvectors are orthogonal, what is the second eigenvector and eigenvalue of Σ ? Don't forget to normalize the eigenvector.

Answer: The spectral decomposition gives us:

$$\begin{aligned} \Sigma - \lambda_1 u_1 u_1^T &= \begin{pmatrix} 50.5 & 49.5 \\ 49.5 & 50.5 \end{pmatrix} - 100 \frac{1}{\sqrt{2}} \begin{pmatrix} 1 \\ 1 \end{pmatrix} \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \end{pmatrix} \\ &= \begin{pmatrix} 50.5 & 49.5 \\ 49.5 & 50.5 \end{pmatrix} - \frac{1}{2} \begin{pmatrix} 100 & 100 \\ 100 & 100 \end{pmatrix} \\ &= \frac{1}{2} \begin{pmatrix} 1 & -1 \\ -1 & 1 \end{pmatrix} \\ &= 1 \times \frac{1}{\sqrt{2}} \begin{pmatrix} 1 \\ -1 \end{pmatrix} \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & -1 \end{pmatrix} \end{aligned}$$


This implies that $\lambda_2 = 1$ and $u_2 = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 \\ -1 \end{pmatrix}$.

- (c) (5 points) What is the “intrinsic” dimensionality of this dataset (discounting some small amount of variance)? Why?

Answer: Clearly the intrinsic dimensionality is 1, since most of the variance ($\frac{100}{101} = 99\%$) is captured by the first principal component.

- (d) (10 points) If the μ and Σ from above characterize the normal distribution from which the points were generated, sketch the exact orientation/extent of the 2D normal in the XY plane. Use the contours corresponding to one standard deviation along each principal axis for your sketch.

Answer: Your sketch should look like this:



2. (15 points) Consider the 3-way contingency table for \mathbf{x} , \mathbf{y} , \mathbf{z} :

	$\mathbf{z}=\mathbf{F}$		$\mathbf{z}=\mathbf{G}$	
	$\mathbf{y}=\mathbf{D}$	$\mathbf{y}=\mathbf{E}$	$\mathbf{y}=\mathbf{D}$	$\mathbf{y}=\mathbf{E}$
$\mathbf{x}=\mathbf{A}$	10	10	10	5
$\mathbf{x}=\mathbf{B}$	15	5	5	20
$\mathbf{x}=\mathbf{C}$	25	10	25	10

(a) (10 points) Compute the χ^2 measure for the correlation between \mathbf{y} and \mathbf{z} .

(b) (5 points) Are they dependent or independent at the 95% confidence level (see the table below for χ^2 values)? Why?

Chi-Square Probabilities: p -values for different Chi-Square values are given for various degrees of freedom df . For example for $df = 5$, a chi-Square value of $\chi^2 = 11.070$ has a p -value of 0.05.

p -value	0.995	0.99	0.975	0.95	0.90	0.10	0.05	0.025	0.01	0.005
df=1	—	—	0.001	0.004	0.016	2.706	3.841	5.024	6.635	7.879
df=2	0.010	0.020	0.051	0.103	0.211	4.605	5.991	7.378	9.210	10.597
df=3	0.072	0.115	0.216	0.352	0.584	6.251	7.815	9.348	11.345	12.838
df=4	0.207	0.297	0.484	0.711	1.064	7.779	9.488	11.143	13.277	14.860
df=5	0.412	0.554	0.831	1.145	1.610	9.236	11.070	12.833	15.086	16.750
df=6	0.676	0.872	1.237	1.635	2.204	10.645	12.592	14.449	16.812	18.548

Answer: Summing out \mathbf{x} , we have the new 2-way contingency table between \mathbf{y} and \mathbf{z} , along with the row/col marginal frequencies:

	$\mathbf{z}=\mathbf{F}$	$\mathbf{z}=\mathbf{G}$	
$\mathbf{y}=\mathbf{D}$	50	40	90
$\mathbf{y}=\mathbf{E}$	25	35	60
	75	75	150

The expected counts in each cell are then given as follows:

	$\mathbf{z}=\mathbf{F}$	$\mathbf{z}=\mathbf{G}$
$\mathbf{y}=\mathbf{D}$	$(90 \cdot 75) / 150 = 45$	$(90 \cdot 75) / 150 = 45$
$\mathbf{y}=\mathbf{E}$	$(60 \cdot 75) / 150 = 30$	$(60 \cdot 75) / 150 = 30$

Subtracting the expected and observed values, and squaring them, we get:

	z=F	z=G
y=D	$5^2=25$	$-5^2=25$
y=E	$-5^2=25$	$5^2=25$

Dividing by the expected counts, gives:

	z=F	z=G
y=A	0.56	0.56
y=B	0.83	0.83

Finally, summing all these values we obtain $\chi^2 = 0.56 + 0.56 + 0.83 + 0.83 = 2.78$.

Since there is only one degree of freedom, we find that the chi-square value is the left of the critical value, namely 3.841, which has a *p-value* of 0.05. Thus we cannot reject the null hypothesis, and we conclude that the two variables are independent.

3. (20 points) Assume that a unit hypercube is given as $[0, 1]^d$, i.e., the domain is $[0, 1]$ in each dimension.

The main diagonal in the hypercube is defined as the vector from $(\mathbf{0}, 0) = (\overbrace{0, \dots, 0}^{d-1}, 0)$ to $(\mathbf{1}, 1) = (\overbrace{1, \dots, 1}^{d-1}, 1)$. For example, when $d = 2$, the main diagonal goes from $(0, 0)$ to $(1, 1)$. On the other hand, the main anti-diagonal is defined as the vector from $(\mathbf{1}, 0) = (\overbrace{1, \dots, 1}^{d-1}, 0)$ to $(\mathbf{0}, 1) = (\overbrace{0, \dots, 0}^{d-1}, 1)$. For example, for $d = 2$, the anti-diagonal is from $(1, 0)$ to $(0, 1)$.

- (a) (10 points) Sketch the diagonal and anti-diagonal in $d = 3$ dimensions, and compute the angle between them.

Answer: The main diagonal is $(1, 1, 1)$ and the anti-diagonal is $(0, 0, 1) - (1, 1, 0) = (-1, -1, 1)$. The angle is therefore: $\cos \theta = \frac{1}{\sqrt{3} \times \sqrt{3}} = 1/3$, which implies $\theta = 70.53^\circ$.

- (b) (10 points) What happens to the angle between the main diagonal and anti-diagonal as $d \rightarrow \infty$. First compute a general expression for the d dimensions, and then take the limit as $d \rightarrow \infty$.

Answer: The main diagonal is $(1, 1, 1)$ and the anti-diagonal is $(\overbrace{-1, \dots, -1}^{d-1}, 1)$.

The angle is therefore: $\cos \theta = -(d-2)/d$.

As $d \rightarrow \infty$, $\cos(\theta) \rightarrow -1 + 2/d = -1$, which implies $\theta = 180^\circ$ or $\theta = 0^\circ$. In other words the diagonal and anti-diagonal are parallel!

4. (15 points) Consider the dataset below, which shows the quantity of each items bought by a customer.

tid	itemset with item quantity
1	2A, 1B, 1C
2	3A, 2B
3	2A, 2B, 1C

Using $\text{minsup} = 2$, find all frequent quantitative itemsets, i.e., frequent itemsets where quantity must be explicitly considered. For example, the frequency of A is 3, the frequency of 2A is 3 (since all three customers buy at least 2 A's), but the frequency of 3A is only 1. You may use/adapt any itemset mining method of your choice.

Answer: The level 1 itemsets A(3), B(3), C(2), all are frequent.

Next level 2: AA(3), AB(3), AC(2), BB(2), BC(2), CC(0), only CC or 2C is not frequent.

Next level 3: AAA(1), AAB(3), AAC(2), ABB(2), ABC(2), BBB(0), BBC(1). The only frequent one are AAB, AAC, ABB, ABC

Final level 4: AABB(2), AABC(2).

5. (10 points) Consider the dataset shown below:

	A	B	Class
x_1	3.5	4	H
x_2	2	4	H
x_3	9.1	4.5	L
x_4	2	6	H
x_5	1.5	7	H
x_6	7	6.5	H
x_7	2.1	2.5	L
x_8	8	4	L

Let us make an “oblique” split, instead of an axis parallel split, given as follows: $A - B \leq 3$. Compute the Information Gain of this oblique split based on Gini Index.

Answer: The gini index for the whole dataset is:

$$1 - (5/8)^2 - (3/8)^2 = 1 - (0.625)^2 - (0.375)^2 = 1 - 0.39 - 0.14 = 0.47.$$

Based on the oblique split, we have $D_Y = \{x_1, x_2, x_4, x_5, x_6, x_7\}$ with $P_H = 5/6$ and $P_L = 1/6$.

For D_N we have $P_H = 0/2 = 0$ and $P_L = 2/2 = 1$.

The gini for D_Y is therefore:


$$1 - (0.83)^2 - (0.17)^2 = 1 - (0.69 + 0.03) = 1 - 0.72 = 0.28.$$

And for D_N it is $1 - 0^2 - 1^2 = 1 - 1 = 0$.

The weighted gini of the split is: $\frac{6}{8}0.28 + \frac{2}{8}0 = 0.21$.

Thus the Gain is $0.47 - 0.21 = 0.26$.

6. (10 points) Consider the dataset shown below:



Define the L_∞ norm, between two points $\mathbf{a} = (a_1, a_2)$ and $\mathbf{b} = (b_1, b_2)$ as follows: $L_\infty(\mathbf{a}, \mathbf{b}) = \max\{|a_1 - b_1|, |a_2 - b_2|\}$. Starting with $\mu_1 = E$ and $\mu_2 = L$, show the clusters after assigning each point to the closest cluster, using the L_∞ distance in the K-means method. In case of ties, assign points to the alphabetically lower center.

Answer: It is clear that A, B, C, D, and E all belong to E. Note that F, G, H, I, J, K all have the same L_∞ distance to both E and L, thus they go to E. The only points that belong to L are: L, M, N.

7. (Bonus: 10 points) Draw a sketch of the 4D hypersphere.